

Automatic Language-Based Verification of Differential Privacy

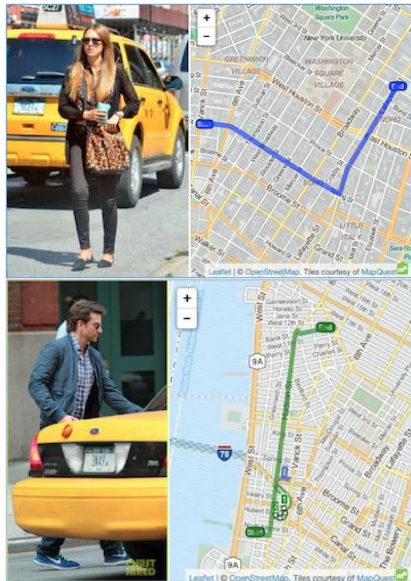
by

Chiké Abuah

Goal: reveal trends in the population, without revealing information about individuals

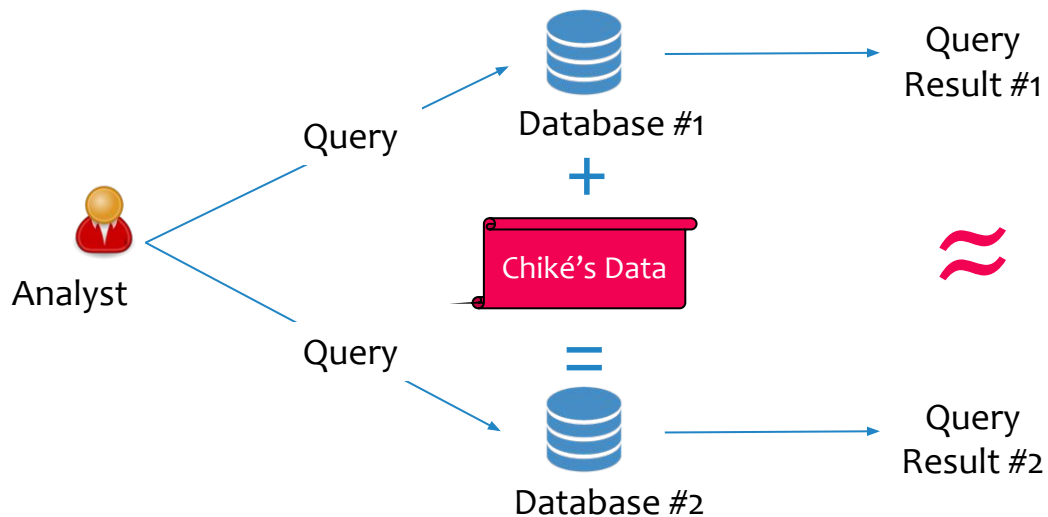
Why Do We Need Differential Privacy?

- Ad-hoc privacy techniques (e.g. anonymization) can't guarantee privacy
- Subject to **re-identification** attacks
 - Netflix recommender algorithm prize (Narayanan et al.)
 - NYC taxi data (Anthony Tocker)



Differential Privacy is a Definition of Privacy

- Adversary can't tell from analysis result whether or not Chiké participated
- Guarantee parameterized by ϵ (the privacy budget)



Satisfying Differential Privacy

To satisfy differential privacy, **add noise to the query result**

$$F(x) = f(x) + \text{Noise}$$

How much noise? **It depends on f**

How many trips were taken in
New York last year?

Low sensitivity

If Chiké lives in Vermont, return
100000; else return 0

High sensitivity

Why do we need Verification?

Algorithm 1 $\mathcal{A}_{\text{Noise-GD}}$: Differentially Private Gradient Descent

Input: Data set: $\mathcal{D} = \{d_1, \dots, d_n\}$, loss function ℓ (with Lipschitz constant L), privacy parameters (ϵ, δ) , convex set \mathcal{C} , and the learning rate function $\eta: [n^2] \rightarrow \mathbb{R}$.

- 1: Set noise variance $\sigma^2 \leftarrow O\left(\frac{L^2 n^2 \log(n/\delta) \log(1/\delta)}{\epsilon^2}\right)$.
 - 2: $\tilde{\theta}_1$: Choose any point from \mathcal{C} .
 - 3: **for** $t = 1$ to $n^2 - 1$ **do**
 - 4: Pick $d \sim_u \mathcal{D}$ with replacement.
 - 5: $\tilde{\theta}_{t+1} = \Pi_{\mathcal{C}}\left(\tilde{\theta}_t - \eta(t) \left[n \nabla \ell(\tilde{\theta}_t; d) + b_t\right]\right)$, $b_t \sim \mathcal{N}(0, \mathbb{I}_p \sigma^2)$.
 - 6: Output $\theta^{\text{priv}} = \tilde{\theta}_{n^2}$.
-

How do we know it's the right amount of noise? *Manual proof*

Incorrect differentially private algorithms *don't crash*
They *silently violate your privacy*

Privacy by (programming language) Design!

- **Accurate**
- **Accessible to non-experts**
- **Automatic**
- **Ubiquitous**
- **Context Matters: Support all manner of programs, systems, software, languages, modes of operation, etc.**

Program Analysis Techniques

Static Analysis

- Control Flow Analysis
- Data Flow Analysis
- Abstract Interpretation
- Type Systems
- Effect Systems
- Model Checking

Dynamic Analysis

- Testing
- Monitoring (Runtime Verification)
- Program Slicing

Talk Summary: Projects

- **Duet**: PL, linear types for (ϵ, δ)
- **DDuo**: dynamic analysis system
- **Solo**: static analysis system

Big Picture Practical Significance

- In practice differential privacy can be difficult to analyze. This work proposes several techniques for **automatic analysis** and **convenient implementation** of differential privacy.
- Ideally, little or no domain knowledge should be required from programmers.
- Privacy violations are **silent**, making it an **especially important field in which to apply verification techniques**.

Threat Model

- Assumes an “honest but fallible” programmer.
- Does not address side-channels.
- Terminated programs can be rerun safely (while consuming the privacy budget).

DUET:

An Expressive Language for Statically
Verifying Differential Privacy

(Published at OOPSLA 2019)

*****ACM SIGPLAN Distinguished Paper Award Winner*****

Duet is *the first linear typed language to support verification of advanced variants of Differential Privacy*

Duet automatically proves that:

```
noisy-gradient-descent( $X, y, k, \epsilon, \delta$ )  $\triangleq$   
  let  $\theta_0 = \text{zeros}(\text{cols } X_1)$  in  
  loop[ $\delta$ ]  $k$  on  $\theta_0 \langle X_1, y \rangle \{t, \theta \Rightarrow$   
     $g_p \leftarrow \text{mgauss}[\frac{1}{m}, \epsilon, \delta] \langle X, y \rangle \{\text{gradient } \theta \ X \ y\} ;$   
    return  $\theta - g_p$  }
```

satisfies

$$(2\epsilon\sqrt{2k\log(1/\delta)}, k\delta + \delta)$$

differential privacy

Duet's Contributions

- Support for **(ϵ, δ) -DP and other variants** in a linear type system
- Expressive **matrix types** and API for machine learning
- Ability to **mix privacy variants** within a single program
- Advanced privacy variants are more accurate in most machine learning scenarios

Duet gets its name...

Duet is a co-design of two distinct, mutually embedded type systems:

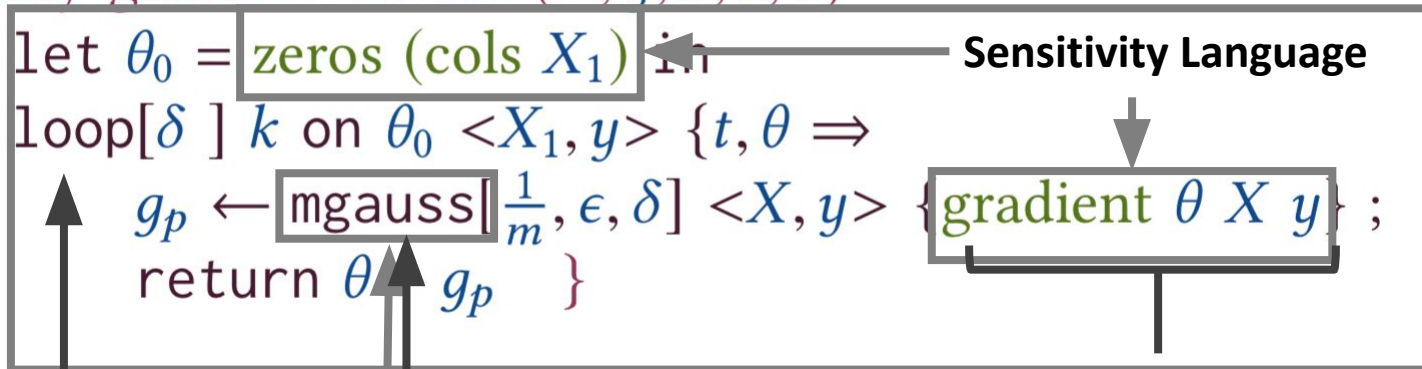
- one for **sensitivity** which leverages linear typing *with scaling* a la Fuzz, and
- one for **privacy** which leverages linear typing *without scaling* and is novel in this work

However, to the user, it's all just one language!



Duet's Two Languages

noisy-gradient-descent($X, y, k, \epsilon, \delta$) \triangleq



Interface between languages:
Gaussian mechanism
Privacy language

Gradient says how to
 improve model
 Sensitivity: $1/m$

Compose many iterations
 with advanced composition

Linear Types for Sensitivity

$x + x$

2-sensitive in x

$$\underbrace{\{x :_2 \mathbb{R}\}} \vdash x + x : \mathbb{R}$$

Context encodes
sensitivity of x

$$\overline{\{x :_1 \tau\} \vdash x : \tau}$$

$$\frac{\Gamma_1 \vdash e_1 : \mathbb{R} \quad \Gamma_2 \vdash e_2 : \mathbb{R}}{\underbrace{\Gamma_1 + \Gamma_2} \vdash e_1 + e_2 : \mathbb{R}}$$

Add up sensitivities
for each variable

Sensitivity in Functions

$$\emptyset \vdash \lambda x : \mathbb{R} \Rightarrow x + x : \mathbb{R} \multimap_2 \mathbb{R}$$

2-sensitive function in x

Type of a 2-sensitive function

Scaling Sensitivity

$$\{y :_4 \mathbb{R}\} \vdash (\lambda x : \mathbb{R} \Rightarrow x + x) (y + y) : \mathbb{R}$$

2-sensitive function in x

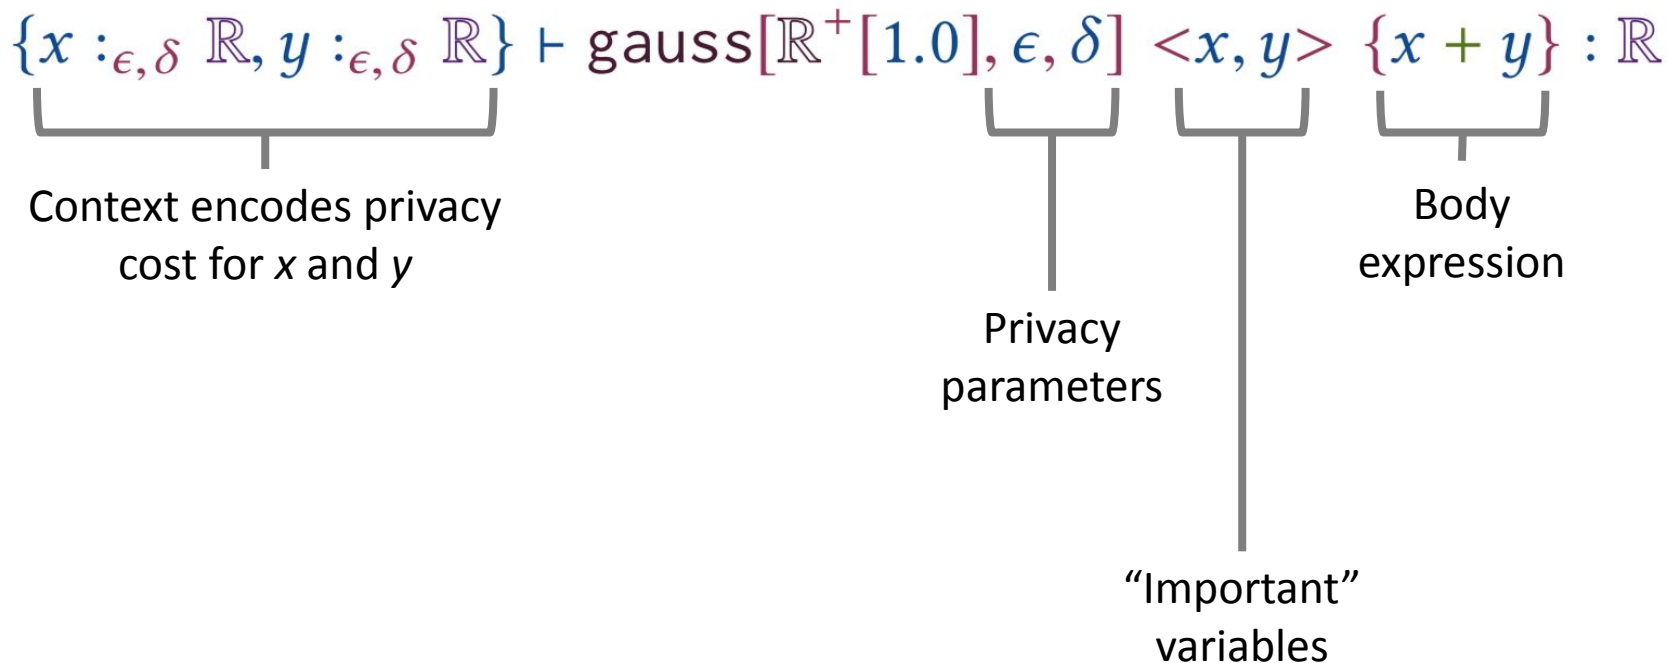
Program is 4-sensitive in y

2-sensitive expression in y

Application rule *scales* context of the argument by the function's sensitivity

$$\frac{\Gamma_1 \vdash e_1 : \tau_1 \multimap_{\circ} \tau_2 \quad \Gamma_2 \vdash e_2 : \tau_1}{\Gamma_1 + s\Gamma_2 \vdash e_1 e_2 : \tau_2}$$

Linear Types for Privacy



Privacy Scaling is **Not Allowed**

$$\{y :_{2\epsilon, 2\delta} \mathbb{R}\} \vdash (p\lambda(x : \mathbb{R}) \Rightarrow \dots) (y + y) : \mathbb{R}$$

Scaling is **allowed** in ϵ -differential privacy (“group privacy”)

Used in *Fuzz* (Reed & Pierce 2010) and *DFuzz* (Gaboardi et al. 2013)

Scaling is **not allowed** in (ϵ, δ) -differential privacy

Duet's Two Languages

\multimap -E

$$\frac{\Gamma_1 \vdash e_1 : \tau_1 \multimap_s \tau_2 \quad \Gamma_2 \vdash e_2 : \tau_1}{\Gamma_1 + s\Gamma_2 \vdash e_1 e_2 : \tau_2}$$

Sensitivity language:
scaling allowed



Privacy language:
no scaling allowed

\multimap^* -E

$$\Gamma \vdash e : (\tau_1 @ p_1, \dots, \tau_n @ p_n) \multimap^* \tau$$

$$\lceil \Gamma_1 \rceil^1 \vdash e_1 : \tau_1 \quad \dots \quad \lceil \Gamma_n \rceil^1 \vdash e_n : \tau_n$$

$$\lceil \Gamma \rceil^\infty + \lceil \Gamma_1 \rceil^{p_1} + \dots + \lceil \Gamma_n \rceil^{p_n} \vdash e(e_1, \dots, e_n) : \tau$$

Theorem (Fundamental Property/Metric Preservation: Sensitivity)

$$\Gamma \vdash e : \tau ; \Sigma \quad , \quad (\gamma_1, \gamma_2) \in \mathcal{G}_{\Sigma'}[\Gamma] \Rightarrow (\gamma_1 \vdash e, \gamma_2 \vdash e) \in \mathcal{E}_{\Sigma'.\Sigma}[\Sigma'(\tau)]$$

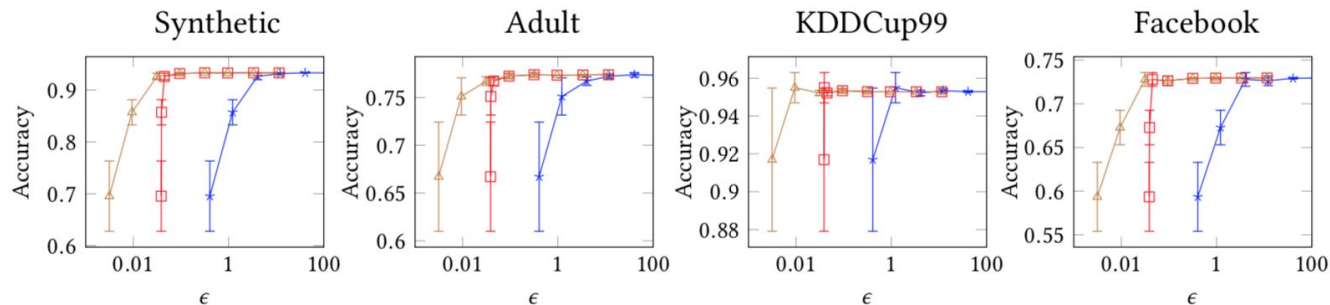
$$\boxed{\begin{array}{l} (v, v) \in \mathcal{V}_s[\tau] \\ (e, e) \in \mathcal{E}_s[\tau] \\ (\gamma, \gamma) \in \mathcal{G}_\Sigma[\Gamma] \end{array}}$$

$$\begin{aligned} (r_1, r_2) \in \mathcal{V}_s[\mathbb{R}] &\stackrel{\Delta}{\iff} |r_1 - r_2| \leq s \\ (\gamma_1 \vdash e_1, \gamma_2 \vdash e_2) \in \mathcal{E}_s[\tau] &\stackrel{\Delta}{\iff} \forall v_1, v_2, \\ &\quad \gamma_1 \vdash e_1 \Downarrow v_1 \wedge \gamma_2 \vdash e_2 \Downarrow v_2 \\ &\quad \implies (v_1, v_2) \in \mathcal{V}_s[\tau] \\ (\gamma_1, \gamma_2) \in \mathcal{G}_\Sigma[\Gamma] &\stackrel{\Delta}{\iff} \forall x \in \text{dom}(\Gamma). (\gamma_1(x), \gamma_2(x)) \in \mathcal{V}_{\Sigma(x)}[\Gamma(x)] \end{aligned}$$

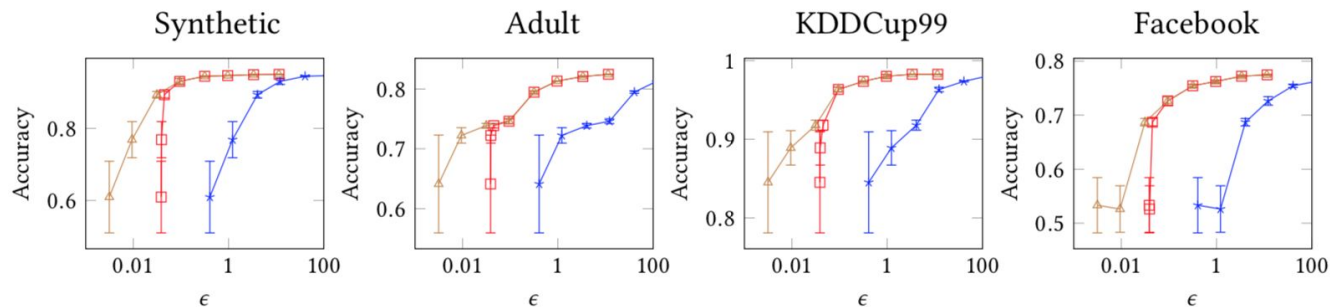
when given related initial configurations and evaluation outputs, then those outputs are related

Empirical Results (Accuracy of Trained Models)

Noisy Gradient Descent



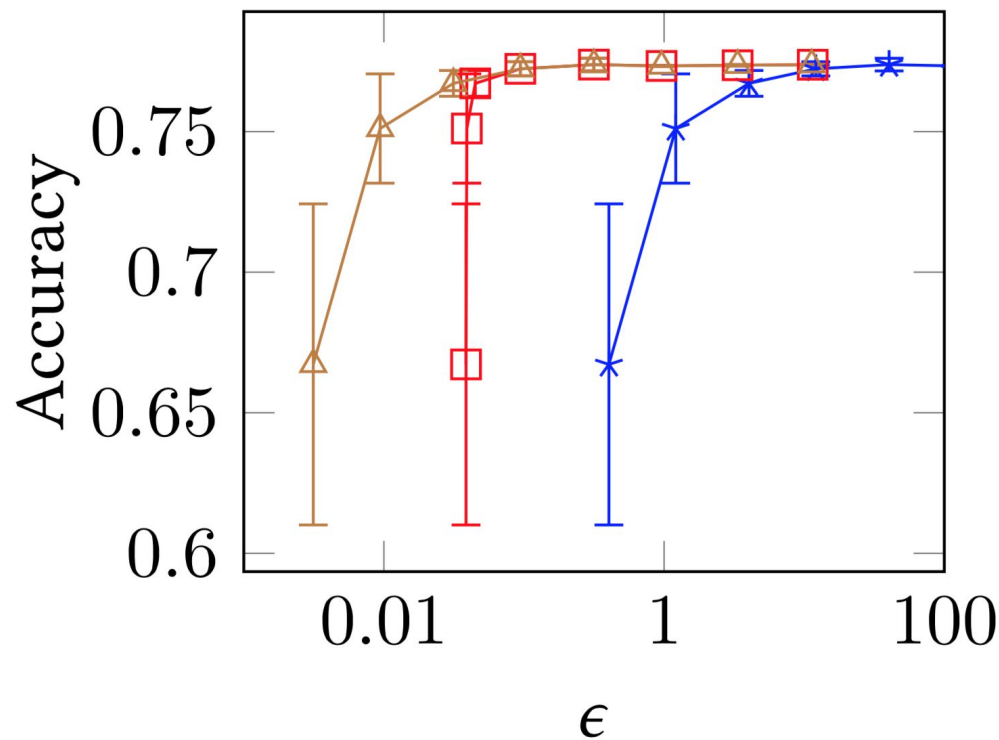
Noisy Frank-Wolfe



Adv. Comp. Rényi DP zCDP

Duet programs
produce
state-of-the-art
accuracy in
private linear
models

Benefit of Differential Privacy Variants



Recent privacy variants yield **better accuracy** at the same level of privacy

—*— Adv. Comp. —□— Rényi DP —△— zCDP

Empirical Results (Typechecker Performance)

Technique	LOC	Time (ms)
Noisy G.D.	23	0.51ms
G.D. + Output Pert.	25	0.39ms
Noisy Frank-Wolfe	31	0.59ms
Minibatching	26	0.51ms
Parallel minibatching	42	0.65ms
Gradient clipping	21	0.40ms
Hyperparameter tuning	125	3.87ms
Adaptive clipping	68	1.01ms
Z-Score normalization	104	1.51ms

Typechecker
computes privacy
cost of complex
programs in
milliseconds

DDUO:

General-Purpose Dynamic Analysis for Differential Privacy

(Published at CSF 2021)

Contributions!

- includes all base types
- general language operations
- various notions of sensitivity
- advanced privacy variants
- generalizes to all future possible arbitrary inputs

DDUO gets its name from...



Product: DDUO makes it easy to automatically enforce privacy

```
def dp_gradient_descent(iterations, alpha, eps):  
    eps_i = eps/iterations  
    theta = np.zeros(X_train.shape[1])  
    noisy_count = duet.renyi_gauss(X_train.shape[0],  $\alpha$  = alpha,  $\epsilon$  = eps)  
    for i in range(iterations):  
        grad_sum = gradient_sum(theta, X_train, y_train, sensitivity)  
        noisy_grad_sum = gaussian_mech_vec(grad_sum, alpha, eps_i)  
        noisy_avg_grad = noisy_grad_sum / noisy_count  
        theta = np.subtract(theta, noisy_avg_grad)  
    return theta
```

- Usable
- For non-experts
- Capable of complex algorithms

DDUO is prototyped in Python

- Does not require the programmer to write any additional type annotations.
- In many cases, DDUO can verify differential privacy for essentially unmodified Python programs.
- easily integrated with popular libraries like Pandas and NumPy.

```
from dduo import pandas as pd
df = pd.read_csv("data.csv")
df

# no change to sensitivity environment
df + 5

# doubles the sensitivity
df + df

( df * 5, df * df)
```

```
with dduo.EpsOdometer() as odo:
    _ = dduo.laplace(df.shape[0],  $\epsilon$  = 1.0)
    _ = dduo.laplace(df.shape[0],  $\epsilon$  = 1.0)
    print(odo)
```

DDUO Overview

- DDUO **data sources** are wrappers around sensitive data.
- DDUO tracks the sensitivity of a value to changes in the program's inputs using a sensitivity environment mapping input *data sources* to *sensitivities*.
- Sensitivity environments of **Sensitive** objects update as operations are applied to them ($f(x) = x + x$).
- DDUO tracks total privacy cost using objects called **privacy odometers**.
- DDUO also allows the analyst to place upper bounds on total privacy cost (i.e. a privacy budget) using **privacy filters**.

```
from dduo import pandas as pd
df = pd.read_csv("data.csv")
df

# no change to sensitivity environment
df + 5

# doubles the sensitivity
df + df

( df * 5, df * df)
```

```
with dduo.EpsOdometer() as odo:
    _ = dduo.laplace(df.shape[0], ε = 1.0)
    _ = dduo.laplace(df.shape[0], ε = 1.0)
    print(odo)
```


Sensitivity Analysis: Object Proxies

```
from dduo import pandas as pd
df = pd.read_csv("data.csv")
df

# no change to sensitivity environment
df + 5

# doubles the sensitivity
df + df

( df * 5, df * df)
```

```
>>> Sensitive(<'DataFrame'>, data.csv ↗ 1,  $L^\infty$ )
```

```
>>> Sensitive(<'DataFrame'>, data.csv ↗ 1,  $L^\infty$ )
```

```
>>> Sensitive(<'DataFrame'>, data.csv ↗ 2,  $L^\infty$ )
```

```
>>> (Sensitive(<'DataFrame'>, data.csv ↗ 5,
 $L^\infty$ ), Sensitive(<'DataFrame'>, data.csv ↗  $\infty$ ,  $L^\infty$ ))
```

Side-effects/Mutation

- Since DDUO attaches sensitivity environments to **values (instead of variables)**, the use of side effects does not affect the soundness of the analysis
- When a program variable is updated to reference a new value, that value's **sensitivity environment remains attached.**
- more capable than traditional type-based static analysis, due to the additional challenges there (e.g. aliasing).

```
total = 0
for i in range(20):
    total = total + df.shape[0]
return total
```

```
>>> Sensitive(<'DataFrame'>,
data.csv ↦ 20, L∞)
```

Conditionals

```
if df.shape[0] == 10:
    return df.shape[0]
else:
    return df.shape[0] * 10000

if dduo.gauss( $\epsilon=1.0$ ,  $\delta=1e-5$ , x) > 5:
    print(dduo.gauss( $\epsilon=1.0$ ,  $\delta=1e-5$ , y))
else:
    print(dduo.gauss( $\epsilon=1000000000000.0$ ,  $\delta=1e-5$ , y))
```

- Branching on sensitive values is disallowed
- We don't want this anyway
- Static Analysis: Take the max!
- Dynamic Analysis: ...?
- **Adaptive** privacy analysis requires use of privacy odometers/filters

Privacy Analysis: Filters and Odometers

```
dduo.laplace(df.shape[0],  $\epsilon=1.0$ )
```

```
with dduo.EpsOdometer() as odo:  
    _ = dduo.laplace(df.shape[0],  $\epsilon = 1.0$ )  
    _ = dduo.laplace(df.shape[0],  $\epsilon = 1.0$ )  
    print(odo)
```

```
with dduo.EdFilter( $\epsilon = 1.0$ ,  $\delta = 10e-6$ ) as odo:  
    print('1:', dduo.gauss(df.shape[0],  $\epsilon=1.0$ ,  $\delta=10e-6$ ))  
    print('2:', dduo.gauss(df.shape[0],  $\epsilon=1.0$ ,  $\delta=10e-6$ ))
```

```
>>> 9.963971319623278
```

```
>>> Odometer_ $\epsilon$ (data.csv ↦ 2.0)
```

```
>>> 1: 10.5627
```

Traceback (most recent call last):

...

dduo.PrivacyFilterException

Loops, Composition, Variants

```
# sequential composition
with dduo.Eps0dometer() as odo:
    for i in range(20):
        dduo.laplace(df.shape[0],  $\epsilon = 1.0$ )
    print(odo)

# advanced composition
with dduo.AdvEd0dometer() as odo:
    for i in range(20):
        dduo.gauss(df.shape[0],  $\epsilon = 0.01$ ,  $\delta = 0.001$ )

# variant mixing
with dduo.Ed0dometer(max_delta = 1e-4) as odo:
    with dduo.RenyiDP(1e-5):
        for x in range(200):
            noisy_count = dduo.renyi_gauss( $\alpha = 10$ ,
                                            $\epsilon=0.2$ , df.shape[0])
    print(odo)
```

- sequential composition

- advanced composition

- variant mixing

Gradient Descent in DDUO

- Bound sensitivity of gradient calculation.
- Add random noise.
- Descend model with noisy gradient.

```
def dp_gradient_descent(iterations, alpha, eps):  
    eps_i = eps/iterations  
    theta = np.zeros(X_train.shape[1])  
    noisy_count = duet.renyi_gauss(X_train.shape[0],  $\alpha$  = alpha,  $\epsilon$  = eps)  
    for i in range(iterations):  
        grad_sum = gradient_sum(theta, X_train, y_train, sensitivity)  
        noisy_grad_sum = gaussian_mech_vec(grad_sum, alpha, eps_i)  
        noisy_avg_grad = noisy_grad_sum / noisy_count  
        theta = np.subtract(theta, noisy_avg_grad)  
    return theta
```

Theorem (Metric Preservation)

If: $\rho_1 \sim_n^\Sigma \rho_2$

And: $\sigma_1 \sim_n^\Sigma \sigma_2$

Then: $\rho_1, \sigma_1, e \sim_n^\Sigma \rho_2, \sigma_2, e$

That is, either $n = 0$ or $n = n' + 1$ and...

If: $n_1 \leq n$

And: $\rho_1 \vdash \begin{bmatrix} \sigma_1, e \end{bmatrix} \Downarrow_{n_1} \begin{bmatrix} \sigma'_1, v_1 \end{bmatrix}$

And: $\rho_2 \vdash \begin{bmatrix} \sigma_2, e \end{bmatrix} \Downarrow_{n_2} \begin{bmatrix} \sigma'_2, v_2 \end{bmatrix}$

Then: $n_1 = n_2$

And: $\sigma'_1 \sim_{n-n_1}^\Sigma \sigma'_2$

And: $v_1 \sim_{n-n_1}^\Sigma v_2$

- the true sensitivity of a program is **guaranteed to be equal to or less than** the sensitivity reported by DDUO's dynamic monitor.
- accurate — even for inputs which **differ entirely from those used** in the dynamic analysis!

Case Studies: Dynamic Enforcement of Privacy

Algorithm	Libraries Used	Baseline	Instrumented Version	Overhead (% increase)
Noisy Gradient Descent	NumPy	5.922s	6.302s	6.42%
Multiplicative Weights (MWEM)	Pandas	0.725s	0.833s	14.90%
Private Naive Bayes Classification	DiffPrivLib	2.155s	2.423s	12.44%
Private Logistic Regression	DiffPrivLib	2.022s	3.161s	56.33%

Solo:

Lightweight Static Analysis for Differential Privacy

(Published at OOPSLA 2022)



Why Solo?

- + **Linear types** provide a strategy rooted in type theory and linear logic for tracking resources throughout the semantics of a core lambda calculus.
- + However, **not commonly available** in mainstream programming languages, and when available are usually **not adequately sophisticated**.



What is Solo?

- + Static type checking for differential privacy in Haskell
- + Doesn't use linear types!
- + Static analysis via (**in tandem** with) Haskell typechecker
- + Regular **Haskell type errors** when privacy is violated
- + **Privacy cost** along with types of functions & values
- + **Encapsulation** of sensitive values via constructor hiding
- + **Automatic** type inference



How does it work?

- + Without linear types, where do we attach sensitivities?
- + Previous dynamic sensitivity analyses have attached sensitivities to values.
- + We embed sensitivities in **base types**—the static equivalent of the dynamic strategy of attaching sensitivities to values.
- + **Type-level parameters** to represent sensitivities symbolically.
- + **Type-level computation** to compute symbolic sensitivity expressions.
- + **Parametric polymorphism** to generalize types over sensitivity parameters.
- + Possible in any language with these e.g. Scala, OCaml and Haskell.



Departure from Linear types

- + **case statements** are *restricted* to have the same sensitivity (or privacy cost) in each case alternative, rather than max. Solved by weaken operation.
- + No verified/typechecked **map or loop**. However trusted primitives are available.
- + No **general recursive datatypes**. However, sensitive collection data types are available.
- + In practice, **none of these are typically barriers** to writing differentially private programs.

Solo gets its name from ...



Basics: Tracking Sensitivity & Privacy

- Data source based approach
- Type-level sensitivity parameters on “sensitive” types
- Parametric effect monad: unifies **privacy effect system** with a monadic-style semantics

```
readDoubleFromIO :: ∀ m o. IO (SDouble m '[ '(o, 1) ])  
  
(<+>) :: SDouble 'Diff senv1  
  -> SDouble 'Diff senv2  
  -> SDouble 'Diff (Plus senv1 senv2)  
  
dbl :: SDouble 'Diff senv -> SDouble 'Diff (Plus senv senv)  
dbl x = x <+> x  
  
simplePrivacyFunction :: SDouble 'Diff '[ '(o, 1) ]  
  -> EpsPrivacyMonad '[ '(o, 2) ] Double  
simplePrivacyFunction x = laplace @2 Proxy $ dbl x
```

Functions



Function types are polymorphic over sensitivity environments

```
-- An s-sensitive function
s_sensitive :: SDouble senv m -> SDouble (ScaleSens senv s) m
-- A 1-sensitive function
one_sensitive :: SDouble senv m -> SDouble senv m
-- map
map :: ∀ m s s1 a b. (∀ s'. a s' -> b (s * s'))
    -> SList m a s1
    -> SList m b (s * s1)
```


Privacy Monad

```
-- values for  $\epsilon$ 
data EpsPrivacyCost = InfEps | EpsCost TLRat
-- privacy environments,  $\epsilon$ -differential privacy
type EpsPrivEnv = [(Source, EpsPrivacyCost)]

return :: a -> EpsPrivacyMonad '[] a
(>=>) :: EpsPrivacyMonad p1 a
  -> (a -> EpsPrivacyMonad p2 b)
  -> EpsPrivacyMonad (EpsSeqComp p1 p2) b
laplace :: Proxy  $\epsilon$ 
  -> SDouble s Diff
  -> EpsPrivacyMonad (TruncateSens  $\epsilon$  s) Double
listLaplace :: Proxy  $\epsilon$ 
  -> L1List (SDouble Diff) s
  -> EpsPrivacyMonad (TruncateSens  $\epsilon$  s) [Double]

addNoiseTwice :: TL.KnownNat (MaxSens s) =>
  SDouble s Diff -> EpsPrivacyMonad '[ '(0,5) ] Double
addNoiseTwice x = do
  y1 <- laplace @2 Proxy x
  y2 <- laplace @3 Proxy x
  return $ y1 + y2
```

- The return operation accepts some value and embeds it in the PrivacyMonad without causing any side-effects.
- The (>=>) (bind) operation allows us to sequence private computations using differential privacy's sequential composition property

Case study: gradient descent

```
-- gradient descent algorithm
gd :: NatS k
    -> NatS  $\epsilon$ 
    -> SMatrix  $\sigma$  LInf m n SDouble
    -> SMatrix  $\sigma$  LInf m 1 SDouble
    -> EpsPrivacyMonad (ScalePriv k (TruncateSens  $\epsilon$   $\sigma$ )) (Matrix 1 n Double)
gd k t xs ys = do
  let m0 = matrix (sn32 @ 1) (sn32 @ n) $ \ i j -> 0
      cxs = mclip xs (natS @ 1)
  let f :: SMatrix  $\sigma_1$  LInf 1 n SDouble
      -> EpsPrivacyMonad (TruncateSens  $\epsilon$   $\sigma_1$ ) (Matrix 1 n Double)
      f = \ $\theta$  -> let g = mlaplace @ $\epsilon$  Proxy (natS @5) $ xgradient  $\theta$  cxs ys
                  in msubM (return  $\theta$ ) g
      z = mloop @(TruncateNat t 1) k (sourceM $ xbp m0) f
  z
```

Proof: similar to DDuo but with type info

Theorem (Metric Preservation)

If: $\gamma_1 \sim \gamma_2 \in \mathcal{G}_n^\Sigma[\Gamma]$

And: $\Gamma \vdash e : \tau @_m^{\Sigma_1}$

Then: $\gamma_1, e \sim \gamma_2, e \in \mathcal{E}_n^{\Sigma \cdot \Sigma_1}[\tau]$

That is, either $n = 0$, or $n = n' + 1$ and...

If: $n'' \leq n$

And: $\gamma_1 \vdash e \Downarrow_{n''} v_1$

Then: $\exists! v_2. \gamma_2 \vdash e \Downarrow_{n''} v_2$

And: $v_1 \sim v_2 \in \mathcal{V}_{n-n''}^{\Sigma \cdot \Sigma_1}[\tau]$



Solo: Case Studies

Solo reference implementation is a (~600 loc) Haskell library

- K-means clustering
- Cumulative Distribution Function
- Gradient Descent
- Multiplicative-Weights Exponential Mechanism (MWEM)

Related Work

Related Work: Statically Typed

Linear Type Systems

- Fuzz, DFuzz (only ϵ -differential privacy)

Indexed Monadic Types

- HOARe² (lacks multi-argument support)


Relational Type Systems

- LightDP (lacks sensitivity analysis)


Type Systems Enriched With Program Logics

- Fuzzi (less support for higher-order/type-checking automation)

Related Work: Dynamically Typed

- PINQ (lacks support for general-purpose programming)
 - ProPer (per-user budget)
 - UniTrax (per-user w/abstract db)
 - Testing Methods (counter-example search)
- 

Related Work: Other Privacy Analysis Software Libraries

- DiffPrivLib (no language-based sensitivity/composition)
 - Google's Privacy Lib (no language-based sensitivity/composition)
 - Ektelo (plans over library of operators, no general PL model)
 - DPella (AST analysis, symbolic interpretation)
- 

Future Work & Conclusions

Open Problem/Future Work

- Property-based testing, sensitivity analysis for large software libraries
- Gradual Differential Privacy

Takeaways!

- Analysis for differential privacy is important: buggy programs silently violate your privacy
- Automated enforcement of privacy can be practical in any context!

<https://github.com/uvm-plaid/duet>

<https://github.com/uvm-plaid/dduo-python>

<https://github.com/uvm-plaid/solo-haskell>

Papers: <https://chikeabuah.github.io/>

Thank You!

Questions?